

# Lab 2 – Exploring Bias, Variance, and Model Evaluation

---

**Time:** 60–90 minutes

**Difficulty:** Beginner → Intermediate

## What you'll learn

- Train/test/validation split (why not just one split?)
  - Understand underfitting vs overfitting through polynomial regression
  - Compare models with different complexity
  - Evaluate using multiple metrics (MSE,  $R^2$ )
  - Visualize learning curves
- 

## Prerequisites

---

- Completion of Lab 1
  - Google Colab account (or local Python with `scikit-learn`, `matplotlib`, `numpy`, `pandas` installed)
- 

## Deliverables

---

1. A Colab notebook with executed cells.
2. Plots showing underfit/overfit behavior.
3. A short reflection (5–8 bullet points) on bias, variance, and generalization.

---

# Step-by-step Instructions

---

## Step 0 – Notebook setup

Open a new notebook in Colab and rename it:

ML-AI-Foundations-Section2-Lab2-YourName

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, learning_curve
from sklearn.metrics import mean_squared_error, r2_score

SEED = 42
np.random.seed(SEED)
```

---

## Step 1 – Generate synthetic data

We'll create noisy quadratic data so learners can experiment with model complexity.

```
# True function: y = 0.5x^2 + x + 2 + noise
X = np.linspace(-3, 3, 100).reshape(-1, 1)
y = 0.5*X**2 + X + 2 + np.random.randn(100, 1)

plt.scatter(X, y, alpha=0.7)
plt.title("Synthetic Data (Quadratic with Noise)")
plt.show()
```

---

## Step 2 – Train/test split

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=SEED  
)
```

---

## Step 3 – Build a helper function for training/evaluation

```
def train_and_evaluate(degree):  
    model = Pipeline([  
        ("poly", PolynomialFeatures(degree=degree, include_bias=False)),  
        ("scaler", StandardScaler()),  
        ("linreg", LinearRegression())  
    ])  
  
    model.fit(X_train, y_train)  
    y_pred_train = model.predict(X_train)  
    y_pred_test = model.predict(X_test)  
  
    train_mse = mean_squared_error(y_train, y_pred_train)  
    test_mse = mean_squared_error(y_test, y_pred_test)  
    r2_train = r2_score(y_train, y_pred_train)  
    r2_test = r2_score(y_test, y_pred_test)  
  
    print(f"Degree {degree}: Train MSE={train_mse:.3f}, Test MSE={test_mse:.3f}")  
  
    # Plot predictions vs data  
    X_range = np.linspace(-3, 3, 100).reshape(-1, 1)  
    y_pred_range = model.predict(X_range)  
    plt.scatter(X_train, y_train, label="Train", alpha=0.6)  
    plt.scatter(X_test, y_test, label="Test", alpha=0.6)  
    plt.plot(X_range, y_pred_range, color="red", linewidth=2, label=f"Model (c  
    plt.legend()  
    plt.title(f"Polynomial Regression (degree {degree})")  
    plt.show()
```

---

## Step 4 – Try different model complexities

```
for d in [1, 2, 10]:  
    train_and_evaluate(d)
```

- **Degree 1 (linear):** underfits → poor fit, low  $R^2$
  - **Degree 2:** good fit → close to true function
  - **Degree 10:** overfits → train perfect, test poor
- 

## Step 5 – Learning curve visualization

Let's see how model performance evolves with more data.

```
train_sizes, train_scores, test_scores = learning_curve(  
    Pipeline([  
        ("poly", PolynomialFeatures(degree=10, include_bias=False)),  
        ("scaler", StandardScaler()),  
        ("linreg", LinearRegression())  
    ]),  
    X, y.ravel(), cv=5,  
    scoring="neg_mean_squared_error",  
    train_sizes=np.linspace(0.1, 1.0, 10),  
    random_state=SEED  
)  
  
train_mean = -train_scores.mean(axis=1)  
test_mean = -test_scores.mean(axis=1)  
  
plt.plot(train_sizes, train_mean, "o-", label="Training error")  
plt.plot(train_sizes, test_mean, "o-", label="Validation error")  
plt.xlabel("Training set size")  
plt.ylabel("MSE")  
plt.title("Learning Curve (Degree=10)")  
plt.legend()
```

```
plt.show()
```

**Expected:** Training error very low, validation error high → variance problem.

---

## Step 6 – Compare metrics (MSE vs R<sup>2</sup>)

Add a markdown cell:

- **MSE (Mean Squared Error):** lower is better; penalizes large errors.
  - **R<sup>2</sup> (coefficient of determination):** closer to 1 = better; <0 means worse than baseline.
- 

## Step 7 – Reflection Questions

Answer in markdown:

1. How did **degree 1 vs 2 vs 10** models differ in fit?
  2. Which one underfit? Which overfit? Why?
  3. What did the learning curve show about bias vs variance?
  4. How does **MSE** compare to **R<sup>2</sup>** as an evaluation metric?
  5. What steps could reduce overfitting? (e.g., regularization, more data, simpler model)
- 

## Step 8 – Save and submit

- Save notebook to Google Drive.
- Export .ipynb and upload as per course instructions.